

Our vision is therefore to extract procedural knowledge (PK) from those documents, and to build a knowledge graph (KG)

according to an ontology. This KG can then be used by different downstream applications to facilitate the access and use of such procedures by human operators.

Examples of such applications could be (KG-empowered) search applications or

intelligent assistant: we expect the users to feel the need to be helped to find a

specific procedure (or a part thereof) and to be guided step-by-step in its ex

ecution, for example by being informed about the action they have to perform

(e.g., turning off a switch), the equipment they may need to use (e.g., wear

ing protective gloves) or the time it may take to perform a specific step (e.g.,

approximately 15 minutes).

In order to fulfill such requirements, the procedural KG extracted from text

should (1) preserve the intended meaning of the original document and (2) con

tain enough information to guide a user in correctly executing the procedure.

The extracted KG should therefore be evaluated, respectively, on the basis of

its quality and usefulness (cf. Section 7). Based on this scenario, we define a

simple ontology and we identify a general-purpose dataset to be used in our

LLM-powered PK extraction and KG building experiments.

Ontology. We reuse existing ontologies when applicable, namely: P-Plan

[14] and K-Hub [42], that address plans and related concepts, FRAPO, one of

the SPAR ontologies<sup>3</sup>, and the Time Ontology<sup>4</sup>, while creating a few new classes

and properties when needed (po:). Specifically, as depicted in Figure 1, a procedure is represented by the class `p-plan:Plan` and is linked to its `p-plan:Steps` (po:hasStep), which are sequentially ordered (property `p-plan:precededBy` and its inverse `khub-proc:nextStep`). Each step is then linked to its `frapo:E`quipment, if any, via the property `frapo:usesEquipment`, and with the action(s) to be performed while executing it (po:hasAction), along with the direct object of the action (po:hasDirectObjectOfAction). The information about the time needed for executing the step is represented by the class `time:TemporalEntity`.

Input procedures. For our experiments, we use as a reference dataset WikiHow, one of the largest online databases of PK, which includes how-to articles on multiple domains. We reuse the JSON dataset built by [52], crawled from the WikiHow website<sup>5</sup>, and focus on atomic procedures, that do not include methods/parts as sub-procedures. Out of each selected procedure, we build an unformatted text (also partially removing punctuation), by concatenating (i) procedure title, representing the overall goal, (ii) general procedure description, (iii) headline of each step, (iv) description of each step. Thus, we also include in the text irrelevant sentences, which are supposed to be discarded during the extraction phase. In total, we used four procedures randomly selected from WikiHow, one as working example for our prompt engineering phase, and three for our method replication and human assessment. Such procedures are

diverse with

respect to both complexity and topic: how to clean a computer monitor (working

example), how to fix a rubbing door, how to make honey glazed parsnips, and how

to plant a bare root tree.